

# Homework 2

STAT6306

## Installing and loading packages

```
packages = c('glmnet', 'leaps')
for(package in packages){
  if(!require(package, character.only=TRUE, quietly=TRUE)){
    install.packages(package, repos='http://cran.us.r-project.org')
    require(package, character.only=TRUE, quietly=TRUE)
  }
}
```

```
## Loaded glmnet 2.0-13
```

## Introduction

A major issue with antiretroviral drugs is the mutation of the virus' genes. Because of its high rate of replication ( $10^9$  to  $10^{10}$  virus per person per day) and error-prone polymerase<sup>1</sup>, HIV can easily develop mutations that alter susceptibility to antiretroviral drugs. The emergence of resistance to one or more antiretroviral drugs is one of the more common reasons for therapeutic failure in the treatment of HIV.

In the paper 'Genotypic predictors of human immunodeficiency virus type 1 drug resistance'<sup>2</sup>, a sample of in vitro<sup>3</sup> HIV viruses were grown and exposed to a particular antiretroviral therapy. The susceptibility of the virus to treatment and the number of genetic mutations of each virus were recorded.

## Question 1

```
load("hiv.rda")

X = hiv.train$x
Y = hiv.train$y

geneLabels = colnames(X)
```

(a)

What are  $n$  and  $p$  in this problem? What are the features in this problem? What are the observations? What is the supervisor? **Note:** Attempt to answer this question before moving on to the rest of the questions.

**#SOLUTION**

---

<sup>1</sup>An enzyme that 'stitches' back together DNA or RNA after replication

<sup>2</sup>The entire paper is on the website. Try to see what you can get out of it if you have the time.

<sup>3</sup>Latin for 'in glass', sometimes known colloquially as a test tube

## SOLUTION

### Question 2

Consider the feature matrix  $\mathbb{X}$ . It is composed of 0's and 1's, with a 1 indicating a mutation in a particular gene. Look at the output for the following chunk of code.

```
table(X)
```

```
## X  
##      0      1  
## 135589 10843
```


What results do you see? What does this indicate?

## SOLUTION

### Question 3

The supervisor is the log transformed susceptibility of a virus to the considered treatment, with large values indicating the virus is relatively more resistant (that is, not susceptible). Run

```
hist(Y)
```



What plot did you just create? What does this indicate?

## SOLUTION

### Question 4

We may have (at least) two goals with a data set such as this:

- inference: can we find some genes whose mutation seems to be most related to viral susceptibility?
- prediction: can we make a model that would predict whether this therapy would be efficacious, given a virus with a set of genetic mutations

(a)

Try to find the best subset solution for this problem. Discuss any problems or findings you discover. In particular, how many possible models are there?

## SOLUTION

(b) Inference

(i)

Find the selected model for:

- forward selection using BIC as the criterion
- lasso

```
#Forward selection  
#SOLUTION
```

```
#lasso: Note that you'll want to consider whether to standardize the feature matrix..  
#SOLUTION
```

(ii)

Compare the selected models for each of the above methods. Which genes are selected by all the methods?

```
#SOLUTION
```

(iii)

At the genes selected by the lasso, which gene mutation sites are associated with a decrease in viral susceptibility to this particular drug? Hint: Consider the signs of the coefficients. What gene site has the largest estimated effect (positive or negative)?

```
#SOLUTION
```

(c) Prediction

(i) Ridge regression

Now that are looking at prediction, we can use ridge regression (which only addresses prediction). Using the package glmnet, plot the CV curve over the grid of  $\lambda$  values and indicate the minimum, and finally report

the CV estimate of the prediction risk for  $\hat{\beta}_{\text{ridge}, \hat{\lambda}}$

**Note:** There is no need to report the  $p$  coefficient estimates from the ridge solution. Also, glmnet has a grid problem. Make two plots, one that shows the problem and one that shows it being corrected.

*#SOLUTION*

## (ii) Prediction on a test set

Now, let's look at some predictions made by these methods. Use the following for the test set:

```
X_0 = hiv.test$x  
Y_0 = hiv.test$y
```

Find an estimate of the risk using the test observations for

- forward selection using  $C_p$  as the criterion
- ridge
- lasso

*#SOLUTION*

## (d)

**Challenge** Suppose we didn't have access to any test data. How could you provide a reasonable estimate of the risk for each of the 3 procedures? What are the pros and cons of your proposal?

*#SOLUTION*

## Question 5

Building on the gradient descent code from lecture, implement (batch or stochastic) gradient descent to find the ridge regression solution (via the Lagrangian formulation) for the HIV data at  $\lambda = 1$ . Compare your solution to the solution found by using the svd and verify that they are (approximately) equal.

Note that the form for the gradients we derived in class need to be augmented by the derivative of the penalty term.

## Additional challenge problems:

I don't want to overwhelm you with homework problems. However, there are additional topics that are relevant for an interested student. You don't need to do these/turn them in.

### LARS vs. forward selection

The LARS algorithm is quite similar to forward selection. Run LARS using the option `forward.stagewise` and compare it to forward selection using Mallows's  $C_p$ .

### GIC-based tuning parameter selection

Try and use a GIC-based method instead of K-fold CV for finding  $\hat{\lambda}$  for the lasso using the HIV data.